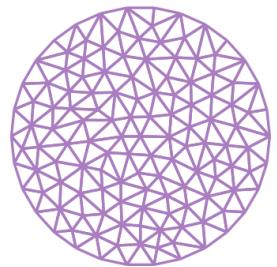
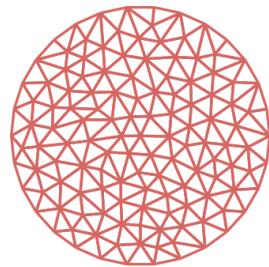
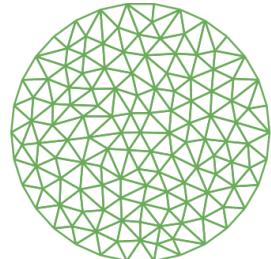


FEM in Julia

An overview of the package
landscape for FEM in Julia



Kristoffer Carlsson

Division of Material and Computational Mechanics
Chalmers University of Technology

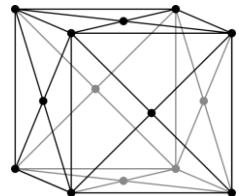
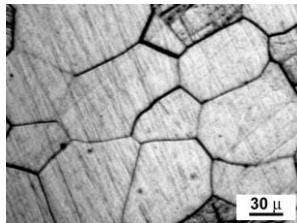
 @KristofferC

kristoffer.carlsson@chalmers.se



CHALMERS

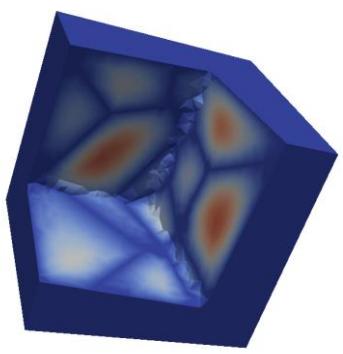
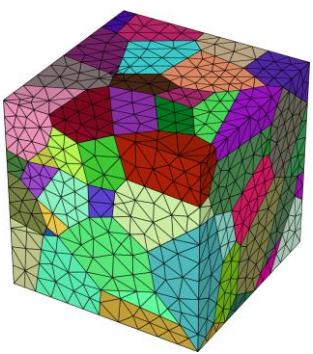
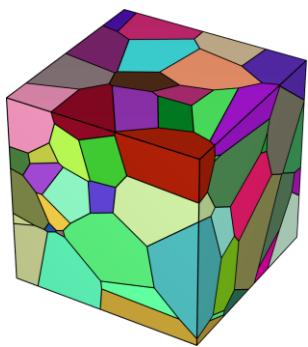
Motivation



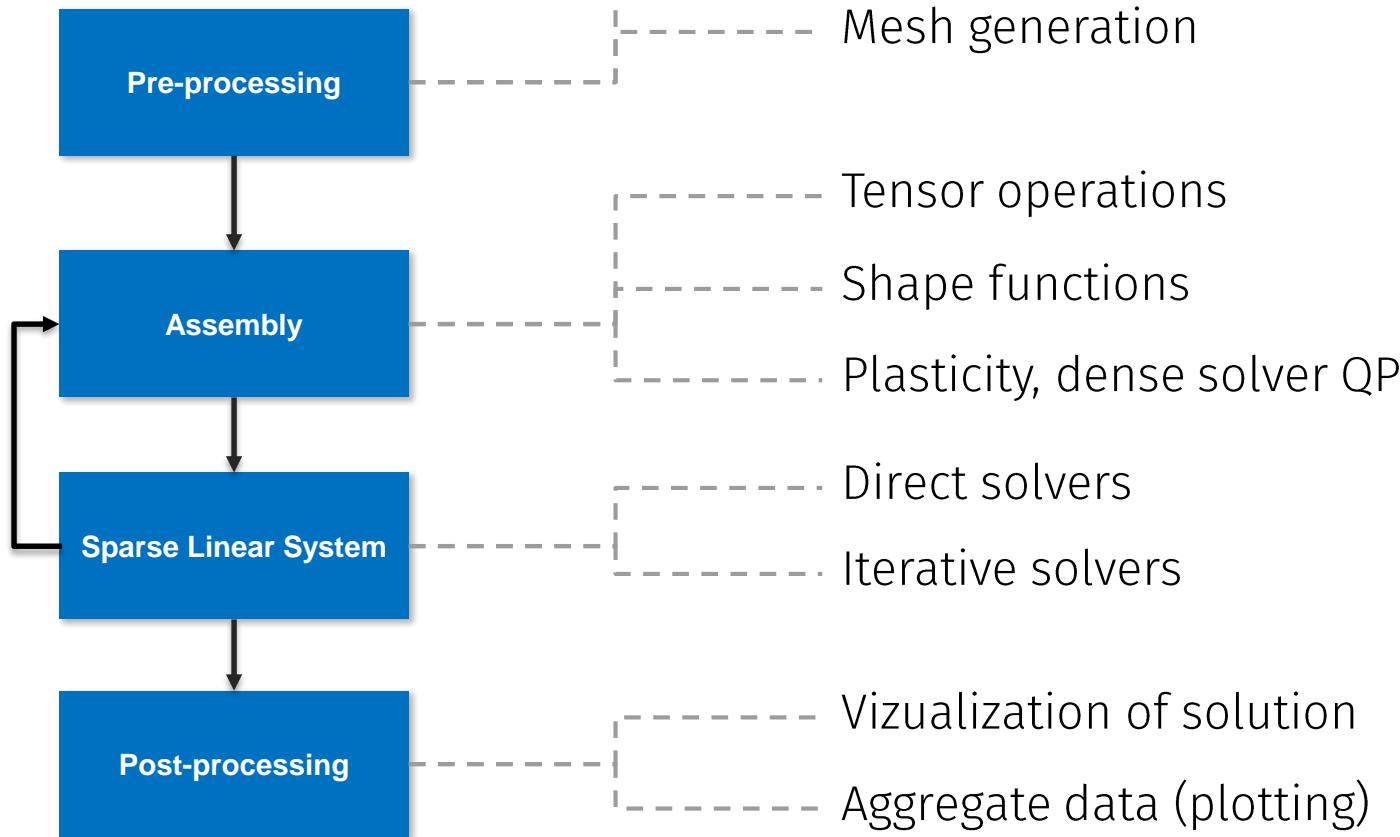
macro $l \approx 10^{-2}\text{m}$ meso $l \approx 10^{-6}\text{m}$ micro $l \approx 10^{-9}\text{m}$

- Macroscopic homogeneous response is a volume average of heterogeneities on a smaller scale
- Material models on macroscopic scale cannot predict new response to changes on the smaller scale

Motivation



FEM Pipeline



Full fledged Julia FEM package?

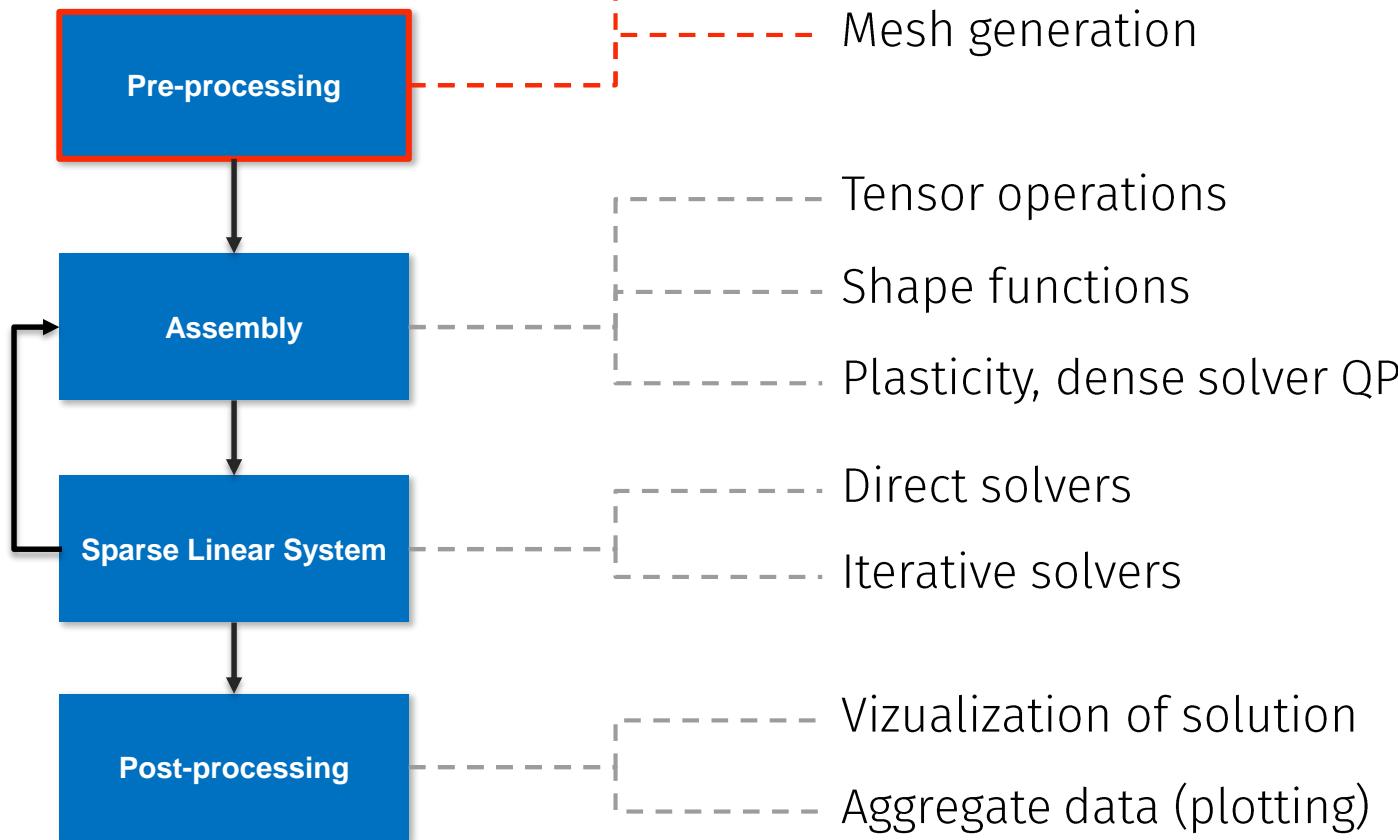
- A full FEM stack is a huge project
 - Deal.II – C++ - 30 000 commits – 20 years
 - DOLFIN – C++ - 21 000 commits – 10 years
- Julia public for \approx 4 years

JuliaFEM.jl - @ahojukka5 et al.

EllipticFEM.jl - @gerhardtulzer

DifferentialEquations.jl - @ChrisRackauckas

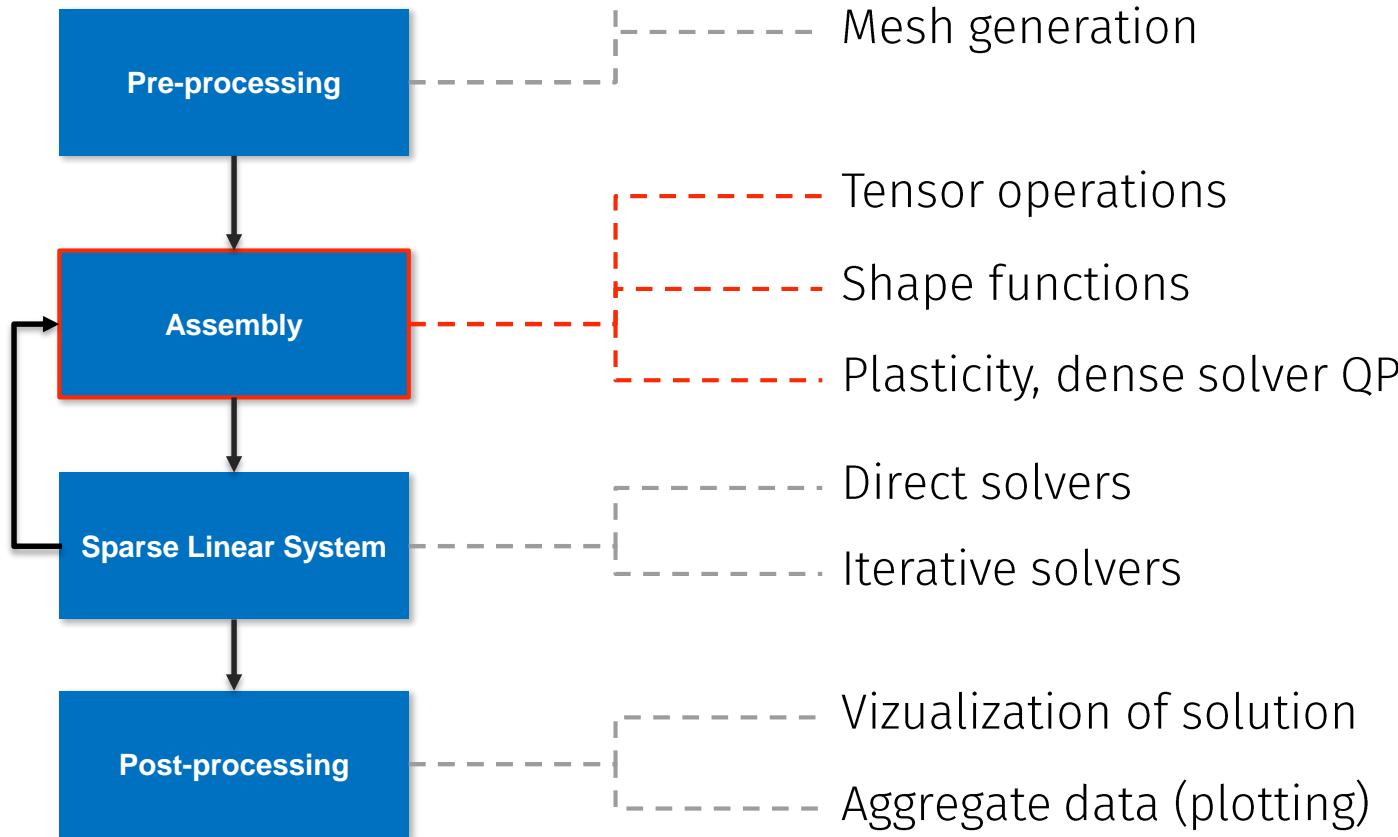
FEM Pipeline



Geometry + Meshing

- Many existing free non-Julia tools for geometry + meshing. GMSH, netgen, tetgen etc.
- Works well, however, need to read the generated mesh files into Julia
- `Meshio.jl` - @SimonDanisch
 - Currently, mostly computer graphics mesh formats
 - I recently added ABAQUS .inp format (in a PR)
 - Call to action!

FEM Pipeline



Tensor operations

- Typically PDE's are given as a collection of tensorial relations
- Working with "naked" Julia arrays as a representation of tensors possible
- This is Julia! No fear of creating new types that maps more directly to our mathematical objects

Tensor operations

- Einsum.jl - @ahwillia

```
using Einsum
A = zeros(5,6,7)
X = randn(5,2)
Y = randn(6,2)
Z = randn(7,2)
@einsum A[i,j,k] = X[i,r]*Y[j,r]*Z[k,r]
```

- TensorOperations.jl – @Jutho

```
using TensorOperations
α=randn()
A=randn(5,5,5,5,5,5)
B=randn(5,5,5)
C=randn(5,5,5)
D=zeros(5,5,5)
@tensor begin
    D[a,b,c] = A[a,e,f,c,f,g]*B[g,b,e] + α*C[c,a,b]
    E[a,b,c] := A[a,e,f,c,f,g]*B[g,b,e] + α*C[c,a,b]
end
```

Tensor operations

- ContMechTensors.jl - @KristofferC
 - Size of tensors known at compile time (1D, 2D, 3D), loop unrolling beneficial
 - Symmetric tensors are common
 - Stack allocation + infix notation → beautiful code

```
using ContMechTensors
A = rand(SymmetricTensor{2, 2})
B = rand(Tensor{4, 2})
x = rand(Vec{2})

A · x # dot product
A ⊗ A # open product
B ⊚ A # double contraction

det(A), inv(A), trace(A)
```

- Maps directly to assembly instructions, no SIMD yet.

Shape functions

- After variational formulation $f = \int_{\Omega_k} v(\boldsymbol{x}) f(\boldsymbol{x}) d\Omega_k$
- Discretization $v \approx v_h = \sum_{i=1}^n V_i \phi_i(\boldsymbol{x})$
- Quadrature $f_i = \sum_{q=1}^{n_{qp}} \phi_i(\boldsymbol{x}_q) f(\boldsymbol{x}_q) w_q$
- Need to evaluate shape functions ϕ at quadrature points defined by $\{\boldsymbol{x}_q, w_q\}_{q=1}^{n_{qp}}$

Shape functions

- JuAFEM.jl – @KristofferC, @fredrikekre
 - Quadrature

```
julia> quadrule = QuadratureRule(:legendre, Dim{2}, RefTetrahedron(), 1);

julia> points(quadrule)
1-element Array{ContMechTensors.Tensor{1,2,Float64,2},1}:
 [0.333333,0.333333]

julia> weights(quadrule)
1-element Array{Float64,1}:
 0.5
```

- FastGaussQuadrature.jl for cubes, tables for tetrahedrons.

Shape functions

- JuAFEM.jl – @KristofferC, @fredrikekre
 - Basis

```
julia> const dim = 2;  
julia> basis_order = 1;  
julia> basis = Lagrange{dim, RefTetrahedron, basis_order}();
```

Shape functions

- JuAFEM.jl – @KristofferC, @fredrikekre
 - FEValues = quadrature + basis

```
julia> fe_values = FEValues(quadrule, basis);  
  
julia> ele_coords = [Vec{2}((0.0, 0.0)),  
                      Vec{2}((1.0, 0.0)),  
                      Vec{2}((1.0, 1.0))]  
  
julia> reinit!(fe_values, ele_coords)  
  
julia> q_point = 1;  
  
julia> node = 2;  
  
julia> shape_value(fe_values, q_point, node)  
0.333333333333333  
  
julia> shape_gradient(fe_values, q_point, node)  
2-element ContMechTensors.Tensor{1,2,Float64,2}:  
 1.0  
 -1.0
```

BlockArrays

- Coupled/mixed problem leads to block like structures

$$R_p(\mathbf{u}, p; \delta p) = 0 \quad \forall \delta p$$

$$R_u(\mathbf{u}, p; \delta u) = 0 \quad \forall \delta \mathbf{u}$$

$$\mathbf{f} = \begin{bmatrix} \mathbf{f}_u \\ \mathbf{f}_p \end{bmatrix} \qquad \mathbf{K} = \left[\begin{array}{c|c} \mathbf{K}_{uu} & \mathbf{K}_{up} \\ \hline \mathbf{K}_{pu} & \mathbf{K}_{pp} \end{array} \right]$$

- Desirable to have a Julia array type that represents this

BlockArrays

- `BlockArrays.jl` – @KristofferC, v0.5 only
- Proposes an `AbstractBlockArray` interface, extension to `AbstractArray`
- Two implementations of block arrays
 - `PseudoBlockArray` – whole matrix stored contiguously
 - `BlockArray` – each block stored contiguously

BlockArrays

```
julia> A = PseudoBlockArray(rand(3,3), [2,1], [2,1])
2×2-blocked 3×3 BlockArrays.PseudoBlockArray{Float64,2,Array{Float64,2}}:
 0.587037 0.443899 | 0.801079
 0.132292 0.196876 | 0.972342
  -----
 0.800054 0.251887 | 0.78099

julia> A[Block(1, 2)]
2×1 Array{Float64,2}:
 0.801079
 0.972342

julia> A[Block(1, 2)] = [0, 0];

julia> A
2×2-blocked 3×3 BlockArrays.PseudoBlockArray{Float64,2,Array{Float64,2}}:
 0.587037 0.443899 | 0.0
 0.132292 0.196876 | 0.0
  -----
 0.800054 0.251887 | 0.78099

julia> nblocks(A)
(2,2)

julia> blocksize(A, 1, 2)
(2,1)

julia> full(A); # returns the "normal" matrix
```

Dense solvers

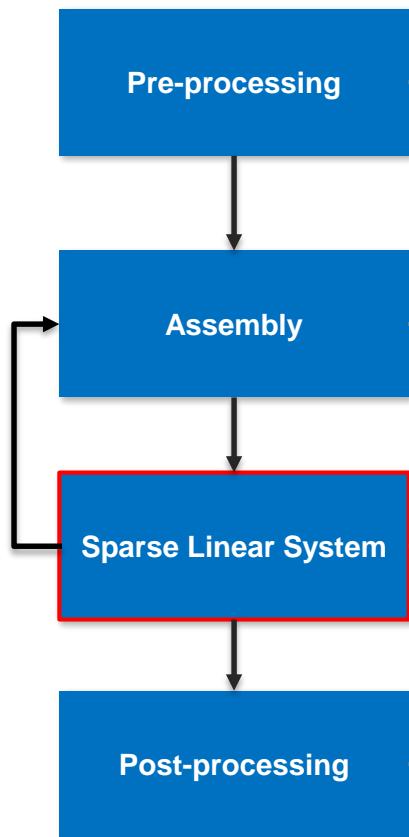
- LAPACK wrapped in Base Julia, 'nuff said?
- Plasticity → maximum dissipation principle → fulfill KKT condition in each quadrature point
- Need to solve small dense nonlinear system, $\approx 10^1$ number of unknowns
- 100 000 elements, 4 quadrature points → 400 000 systems to solve per Newton iteration
- Overhead in solver packages can dominate

```
function solve(f, x0)
    # allocate temporary arrays
    for i in iter
        do_step(f, x)
    end
    return solution
end
```

```
function solve(f, x0, cache=SolveCache(x0))
    # unpack cache (Parameters.jl, @mauro3)
    for i in iter
        do_step(f, x)
    end
    return solution
end
```

- NLsolve.jl, Optim.jl + ForwardDiff.jl = ❤️

FEM Pipeline



Geometry

Mesh generation

Shape functions

Tensor operations

Plasticity, dense solver QP

Direct solvers

Iterative solvers

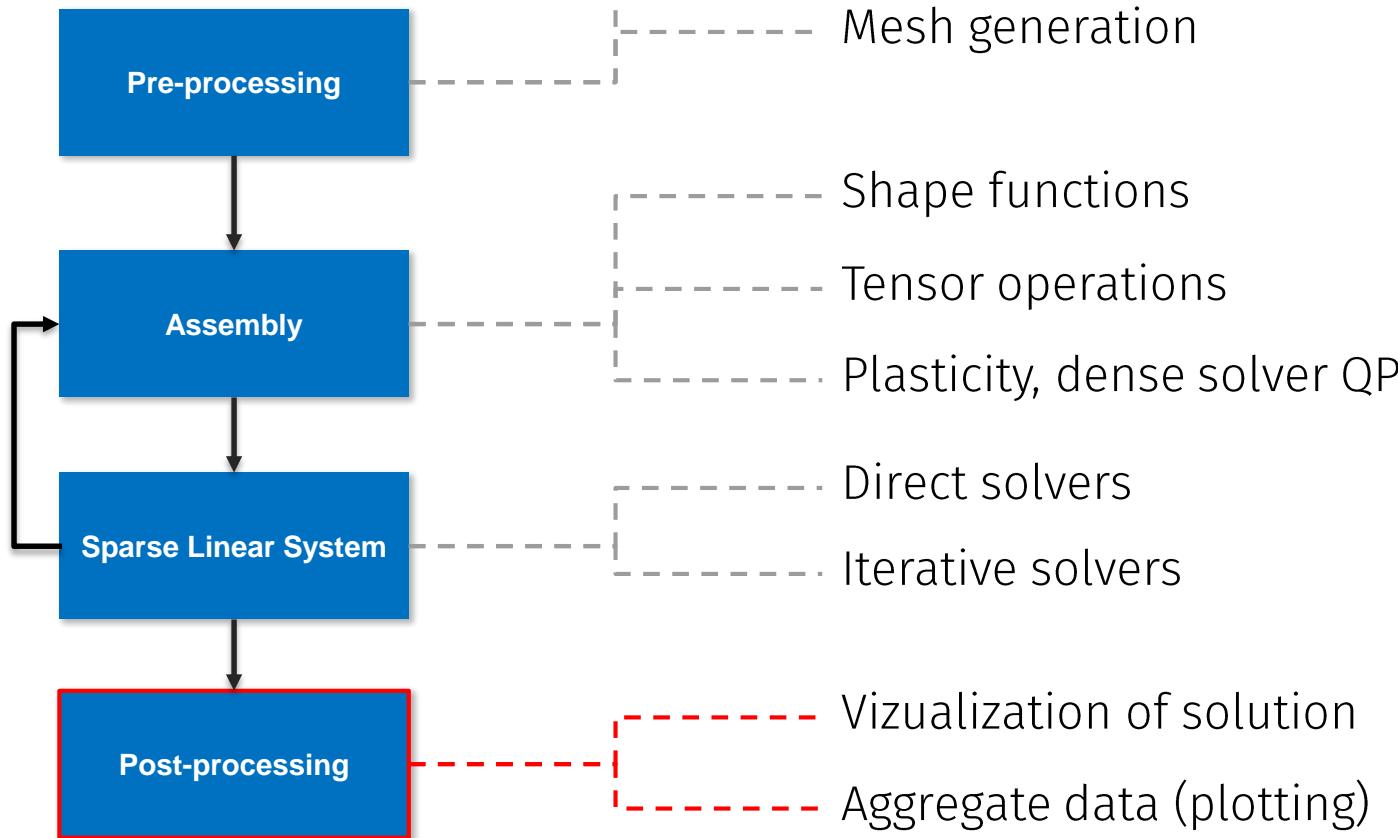
Vizualization of solution

Aggregate data (plotting)

Sparse Linear Solvers

- Direct solvers
 - Cholmod + UMFPACK – Comes with Julia, "\\"
 - Pardiso.jl, Mumps.jl, Wrappers to C libraries
- Iterative solvers
 - IterativeSolvers.jl - @jiahao et al.
 - KrylovMethods.jl - @lruthotto
 - Krylov.jl - @dpo
 - PyAMG.jl - @cortner – PyAMG wrapper with PyCall
 - See Lars Ruthotto: [Iterative Methods for Sparse Linear Systems in Julia: A Quick Overview](#)
- PETSc.jl – Jared Crean – Talk on Friday

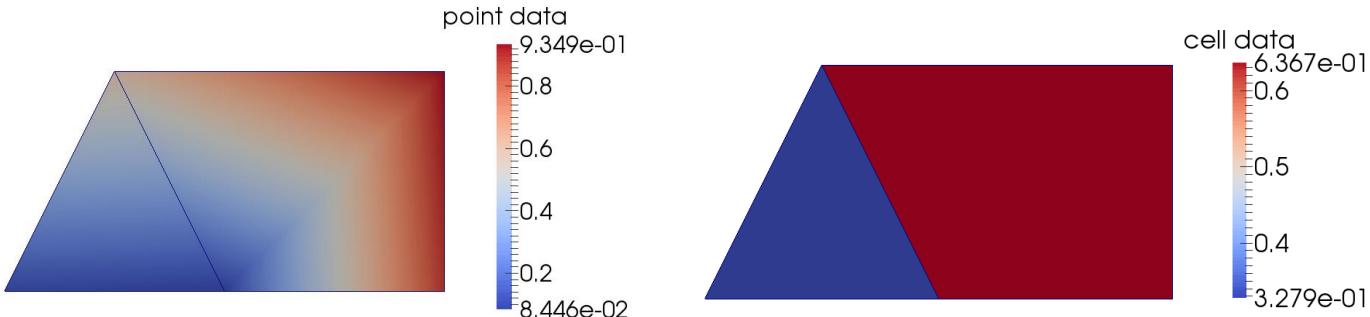
FEM Pipeline



Vizualization of solutions

- Dominating format: .VTK + Paraview/Visit
- WriteVTK.jl - @jipolanco

```
using WriteVTK
points = [0.0 0.5 2.0 1.0 2.0;
          0.0 1.0 0.0 0.0 1.0;
          0.0 0.0 0.0 0.0 0.0]
cells = [MeshCell(VTKCellType.VTK_TRIANGLE, [1, 4, 2]),
         MeshCell(VTKCellType.VTK_QUAD, [2, 4, 3, 5])]
vtkfile = vtk_grid("my_vtk_file", points, cells)
vtk_point_data(vtkfile, rand(5), "point data")
vtk_cell_data(vtkfile, rand(2), "cell data")
vtk_save(vtkfile)
```



Aggregate data (plotting)

- Plotting landscape in Julia used to be quite fragmented
- Thanks to `@tbreloff` and `Plots.jl` this is no longer the case
- `Dataframes.jl + JLD.jl + Plots.jl`
(`PGFPlots.jl` backend for publications)

```
using JLD
using Dataframes
using Plots
pgfplots()

df = load("data.jld")["analysis_data"]
data = extract_data(df)
p = plot(data)
PGFPlots.save("pgfplot.tex", p.o, include_preamble=false)
```

Kristoffer's wishlist

- Make working with stack allocated arrays easier
 - `Array` – powerful but big and heap allocated
 - Goto solution is tuples
 - Generated code explosions, type inference problems, immutability not always desired
 - See issue #11902
- Facilitate not double paying lookup cost when mutating `LinearSlow` elements
 - `K[i,j] += 1.0 → K[i,j] = K[i,j] + 1.0`
 - `K sparse → [i,j]` expensive, double paying
 - See issue #15630 and
<https://github.com/KristofferC/UpdateIndex.jl>

Conclusions

- Package landscape for FEM is developed well enough to do actual work
- State of the art Automatic Differentiation makes local stiffness tangents a breeze to evaluate
- Possible to use Julia for only a part of the FEM stack by calling into Julia from C++ FEM libraries

Thank you!

Slides

https://github.com/KristofferC/JuliaCon_FEM

Acknowledgements

Swedish Research Council



Vetenskapsrådet

Julia Computing

